

Sample Project: Using pyTorch to classify flowers by species.

Similar to many machine learning project, my approach includes:

- image processing
- Training classifier
- Deployment of the classifier.

I didn't train my own model, it will take too long, instead I used transform learning by integrate pretrained models. I'll be using [this dataset](http://www.robots.ox.ac.uk/~vgg/data/flowers/102/index.html) (<http://www.robots.ox.ac.uk/~vgg/data/flowers/102/index.html>) of 102 flower categories.

```
In [ ]: import os
import json
import numpy as np
import pandas as pd
import torch # using pytorch
from torch import optim
from torchsummary import summary # borrow from https://github.com/sksq96/pytorch-summary to show the model configuration
from timeit import default_timer as timer

## Computer Vision functions
from torchvision import transforms, datasets
from torchvision import models
from PIL import Image

## Neural network functions
import torch.nn as nn
import torch.nn.functional as F

## data loading functions
from torch.utils.data import DataLoader
from torch.utils.data.sampler import SubsetRandomSampler

## plotting function
import matplotlib.pyplot as plt
%matplotlib inline

## to display the full results instead of the last one.
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
import sys
```

Load the data

```
In [ ]: batch_size = 32
        IMAGE_SIZE = 224
        USE_GPU = torch.cuda.is_available()
        MODEL_SAVE_FILE = 'XI_Model.pth'

        data_dir = './assets/flower_data'
        train_dir = data_dir + '/train'
        valid_dir = data_dir + '/valid'
```

```
In [ ]: # Transforming for the training and validation sets
# Images need to be processed to the right size, and normalized with the mean
# of the whole data size.
training_transforms = transforms.Compose([
    transforms.RandomRotation(45),
    transforms.RandomResizedCrop(IMAGE_SIZE), # adding randomness
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
                        [0.229, 0.224, 0.225])
])

validation_transforms = transforms.Compose([
    transforms.Resize(IMAGE_SIZE + 32),
    transforms.CenterCrop(IMAGE_SIZE),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
                        [0.229, 0.224, 0.225])
])

# Load the datasets with ImageFolder function with transform function passed
# in
train_dataset = datasets.ImageFolder(train_dir, transform =training_transforms)
valid_dataset = datasets.ImageFolder(valid_dir, transform =validation_transforms)

# Define the dataloaders, which is similar to "generator"
train_dataloader = torch.utils.data.DataLoader(train_dataset, batch_size =batch_size, shuffle =True)
valid_dataloader = torch.utils.data.DataLoader(valid_dataset, batch_size =batch_size)

dataloaders = {'train': train_dataloader,
               'valid': valid_dataloader}

dataset_sizes = {'train': len(train_dataset),
                 'valid': len(valid_dataset)}
```

```
In [ ]: # Loading labels

with open('cat_to_name.json', 'r') as f:
    cat_to_name = json.load(f)

list(cat_to_name.items())[:5]
```

Building and training the classifier

I will use DenseNet model, which I feel should give me a great result.

```
In [ ]: DenseNet = models.densenet161(pretrained=True)
        #print(DenseNet)
        # freeze all pretrained model parameters
        for param in DenseNet.parameters():
            param.requires_grad_(False)
```

```
In [ ]: # My classifier:
        # modify the last layer by adding two fully connected layer and an log softmax
        # layer as output.

        from collections import OrderedDict
        classifier = nn.Sequential(OrderedDict([
            ('fc1', nn.Linear(2208, 1200)),
            ('relu', nn.ReLU()),
            ('fc2', nn.Linear(1200, 102)),
            ('output', nn.LogSoftmax(dim=1))
        ]))
        DenseNet.classifier = classifier
```

```
In [ ]: # Specify Loss Function and Optimizer
        # Criteria NLLLoss which is recommended with Softmax final layer
        criterion = nn.NLLLoss()

        # Set Learning rate and decay Learning rate by a factor of 0.1 every 4 epoches
        optimizer = optim.Adam(DenseNet.classifier.parameters(), lr=0.001)

        from torch.optim import lr_scheduler
        scheduler = lr_scheduler.StepLR(optimizer, step_size=4, gamma=0.1)
```

Install needed package for Google colab

```
In [ ]: !pip install requests  
        !pip install airtable
```

```
In [ ]: def train_model(model, criterion, optimizer, scheduler, n_epoch=60):  
    ''' train_model  
    Model training function.  
  
    Argument:  
        model: model archetecture  
        criterion: loss function  
        optimizer: optimization function  
        scheduler: schedular  
        n_epoch: number of epoches to train on  
    Return: trained model  
    '''  
  
    # if torch.cuda.is_available():  
    model = model.cuda()  
  
    best_model_wts = copy.deepcopy(model.state_dict())  
    best_acc = 0.0  
  
    start_time = timer()  
  
    for epoch in range(n_epoch):  
        print('Epoch {}/{}'.format(epoch, n_epoch-1))  
        print('-' * 10)  
  
        # Each epoch has a training and validation phase  
        for phase in ['train', 'valid']:  
            if phase == 'train':  
                scheduler.step()  
                model.train() # Set model to training mode  
            else:  
                model.eval() # Set model to evaluate mode  
  
            running_loss = 0.0  
            running_corrects = 0  
  
            # Iterate over data.  
            for inputs, labels in dataloaders[phase]:  
                inputs = inputs.cuda()  
                labels = labels.cuda()  
  
                # zero the parameter gradients
```

```

optimizer.zero_grad()

# forward
# track history if only in train
with torch.set_grad_enabled(phase == 'train'):
    outputs = model(inputs)
    _, preds = torch.max(outputs, 1)
    loss = criterion(outputs, labels)

    # backward + optimize only if in training phase
    if phase == 'train':
        loss.backward()
        optimizer.step()

    # statistics
    running_loss += loss.item() * inputs.size(0)
    running_corrects += torch.sum(preds == labels.data)
epoch_loss = running_loss / dataset_sizes[phase]
epoch_acc = running_corrects.double() / dataset_sizes[phase]

print('{} Loss: {:.4f} Acc: {:.4f}'.format(
    phase, epoch_loss, epoch_acc))

# deep copy the model
if phase == 'valid' and epoch_acc > best_acc:
    best_acc = epoch_acc
    best_model_wts = copy.deepcopy(model.state_dict())

print()
time_elapsed = timer() - start_time
print('Training complete in {:.0f}m {:.0f}s'.format(
    time_elapsed // 60, time_elapsed % 60))
print('Best val Acc: {:.4f}'.format(best_acc))

# Load best model weights
model.load_state_dict(best_model_wts)
return model

```

```
In [ ]: model_ft = train_model(DenseNet, criterion=criterion, optimizer=optimizer, scheduler=scheduler, n_epoch=60)
```


Save the checkpoint

```
In [ ]: def save_checkpoint(model, save_path, save_cpu=False):  
    ''' Same checkpoint to local  
    Argument:  
        model: model to save  
        save_path: local path for save model  
        save_cpu: whether to conver model to CPU.  
    Return: None  
    '''  
    if save_cpu:  
        model = model.cpu()  
  
    checkpoint = {  
        'arch': 'DenseNet',  
        'state_dict': model.state_dict()  
    }  
  
    torch.save(checkpoint, save_path)
```

```
In [ ]: save_checkpoint(model_ft, "./DenseNet.pth", save_cpu=True)
```

Loading the checkpoint

At this point it's good to write a function that can load a checkpoint and rebuild the model. That way you can come back to this project and keep working on it without having to retrain the network.

```
In [ ]: def load_model(checkpoint_path):  
    ''' Load trained model from checkpoint  
    Argument: checkpoint path  
    Return: model  
    '''  
  
    from collections import OrderedDict  
    chpt = torch.load(checkpoint_path)  
    model = models.densenet161(pretrained=True)  
    for param in model.parameters():  
        param.requires_grad = False  
  
    #     model.class_to_idx = chpt['class_to_idx']  
  
    # Create the classifier  
    classifier = nn.Sequential(OrderedDict([  
        ('fc1', nn.Linear(2208, 1200)),  
        ('relu', nn.ReLU()),  
        ('fc2', nn.Linear(1200, 102)),  
        ('output', nn.LogSoftmax(dim=1))  
    ]))  
    # Put the classifier on the pretrained network  
    model.classifier = classifier  
  
    model.load_state_dict(chpt['state_dict'], strict=False)  
  
    return model
```

```
In [ ]: load_model("DenseNet1.pth")
```

Deployment for prediction.

```
In [ ]: def process_image(image):  
        ''' Process Image function for prediction  
        Argument: an input image  
        Return: a processed image  
        '''  
  
        image = Image.open(image)  
  
        image_transforms = transforms.Compose([transforms.Resize(256),  
                                              transforms.CenterCrop(224),  
                                              transforms.ToTensor(),  
                                              transforms.Normalize([0.485,0.456,0.406],  
                                                                  [0.229,0.224,0.225])])  
  
        image = image_transforms(image)  
        image = np.array(image)  
        return image
```

To check your work, the function below converts a PyTorch tensor and displays it in the notebook. If your `process_image` function works, running the output through this function should return the original image (except for the cropped out portions).

```
In [ ]: def imshow(image, ax=None, title=None):  
    """Imshow for Tensor."""  
    if ax is None:  
        fig, ax = plt.subplots()  
  
    # PyTorch tensors assume the color channel is the first dimension  
    # but matplotlib assumes is the third dimension  
    image = image.numpy().transpose((1, 2, 0))  
  
    # Undo preprocessing  
    mean = np.array([0.485, 0.456, 0.406])  
    std = np.array([0.229, 0.224, 0.225])  
    image = std * image + mean  
  
    # Image needs to be clipped between 0 and 1 or it looks like noise when displayed  
    image = np.clip(image, 0, 1)  
  
    ax.imshow(image)  
  
    return ax
```

```
In [ ]: def predict(image_path, model, topk=5):  
    ''' Predict the class of an image using a trained model.  
  
    Argument:  
        image_path: image location for classification  
        model: trained model.  
        topk: show top k classes with highest predicted prob.  
    Return:  
        probs: probabilities,  
        classes: predicted classes  
    ...  
    model.eval()  
  
    probs = torch.exp(model(img))  
  
    return print(probs, classes)
```